

# CMSC 125: Operating Systems

- ❑ Instructor: **Joseph Anthony C. Hermocilla**
- ❑ Email: [jchermocilla@up.edu.ph](mailto:jchermocilla@up.edu.ph)
- ❑ Web: <https://jachermocilla.org>



# Resources

Book: <https://pages.cs.wisc.edu/~remzi/OSTEP/>

Slides Template:

<https://pages.cs.wisc.edu/~remzi/OSTEP/Educators-Slides/Youjip/>



# Acknowledgement

- ▣ This lecture slide set was initially developed for Operating System course in Computer Science Dept. at Hanyang University. This lecture slide set is for OSTEP book written by Remzi and Andrea at University of Wisconsin.

# **26. Concurrency: An Introduction**

**Operating System: Three Easy Pieces**

# Thread


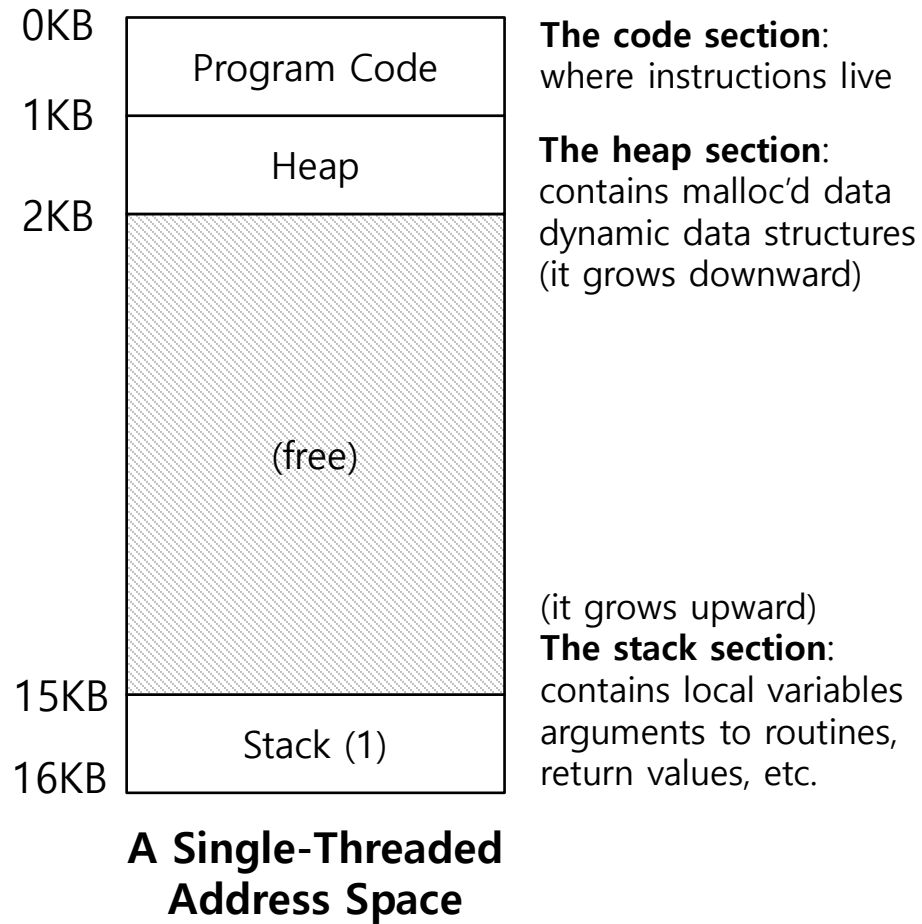
- A new abstraction for a single running process
- Multi-threaded program
  - ◆ A multi-threaded program has more than one path of execution
  - ◆ Multiple PCs (Program Counter)
  - ◆ They **share** the same **address space**

# Context switch between threads

- Each thread has its own program counter and set of registers
  - ◆ One or more **thread control blocks(TCBs)** are needed to store the state of each thread
- When switching from running one (T1) to running the other (T2)
  - ◆ The register state of T1 be saved
  - ◆ The register state of T2 restored
  - ◆ The **address space remains** the same

# The stack of the relevant thread

- There will be **one stack per thread**



# Why Use Threads?

1. Parallelism
  - ◆ Speeds up computations
2. Avoid blocking a process' progress due to slow I/O
  - ◆ Enables the overlap of I/O and computations **within a process**
  - ◆ Ex. Spell-checking using a background thread in a word processor



# Thread Creation Example

```
1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>
4  #include "common.h"
5  #include "common_threads.h"
6
7  void *mythread(void *arg) {
8      printf("%s\n", (char *) arg);
9      return NULL;
10 }
11
12 int
13 main(int argc, char *argv[]) {
14     pthread_t p1, p2;
15     int rc;
16     printf("main: begin\n");
17     Pthread_create(&p1, NULL, mythread, "A");
18     Pthread_create(&p2, NULL, mythread, "B");
19     // join waits for the threads to finish
20     Pthread_join(p1, NULL);
21     Pthread_join(p2, NULL);
22     printf("main: end\n");
23     return 0;
24 }
```

# Thread Creation Example: Possible Execution Path 1

main	Thread 1	Thread2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
waits for T1	runs	
	prints "A"	
	returns	
waits for T2		runs
		prints "B"
		returns
prints "main: end"		

## Thread Creation Example: Possible Execution Path 2

main	Thread 1	Thread2
starts running		
prints "main: begin"		
creates Thread 1		
	runs	
	prints "A"	
	returns	
creates Thread 2		
		runs
		prints "B"
		returns
waits for T1		
<i>returns immediately; T1 is done</i>		
waits for T2		
<i>returns immediately; T2 is done</i>		
prints "main: end"		

## Thread Creation Example: Possible Execution Path 3

main	Thread 1	Thread2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		runs
		prints "B"
		returns
waits for T1	runs	
	prints "A"	
	returns	
waits for T2		
<i>returns immediately; T2 is done</i>		
prints "main: end"		

# Why So Many Possible Execution Paths?

- Depends on the thread scheduler!

# Shared Data Example

- Further complicates things because data inconsistencies can happen

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include "common.h"
4  #include "common_threads.h"
5
6  static volatile int counter = 0;
7
8  // mythread()
9  //
10 // Simply adds 1 to counter repeatedly, in a
11 // No, this is not how you would add 10,000,
12 // a counter, but it shows the problem nicely
13 //
14 void *mythread(void *arg) {
15     printf("%s: begin\n", (char *) arg);
16     int i;
17     for (i = 0; i < 1e7; i++) {
18         counter = counter + 1;
19     }
20     printf("%s: done\n", (char *) arg);
21     return NULL;
22 }
```

**Critical Section**

```
mov 0x8049a1c, %eax
add $0x1, %eax
mov %eax, 0x8049a1c
```

## Shared Data (Cont..)


```
23
24 // main()
25 //
26 // Just launches two threads (pthread_create)
27 // and then waits for them (pthread_join)
28 //
29 int main(int argc, char *argv[]) {
30     pthread_t p1, p2;
31     printf("main: begin (counter = %d)\n", counter);
32     Pthread_create(&p1, NULL, mythread, "A");
33     Pthread_create(&p2, NULL, mythread, "B");
34
35     // join waits for the threads to finish
36     Pthread_join(p1, NULL);
37     Pthread_join(p2, NULL);
38     printf("main: done with both (counter = %d)\n",
39           counter);
40     return 0;
41 }
```

# Race Condition

- The result depends on the execution path
- Example with two threads

- ◆ counter = counter + 1 (initial is 50)
- ◆ We expect the result is 52. However,

Part of thread state



OS	Thread1	Thread2	PC	%eax	counter
	before critical section		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
		mov 0x8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0x8049a1c	113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	50
	mov %eax, 0x8049a1c		113	51	<b>51</b>



## Critical Section

- ❑ A piece of code that **accesses a shared variable** and must not be concurrently executed by more than one thread because executing it might result to a race condition and thus to incorrect results
- ❑ Solution: Need to support **atomicity**(all instruction executions in the critical section will not be preempted) for critical section access (**mutual exclusion**)
  - ◆ Will need **special hardware instructions** and **OS support**

# Locks

- Another abstraction such that a thread will not be able to execute in the critical section unless it holds a lock
- Ensure that any such critical section executes as if it were a single atomic instruction (**execute a series of instructions atomically**).

```
1  lock_t mutex;  
2  . . .  
3  lock(&mutex);  
4  balance = balance + 1;  
5  unlock(&mutex);
```

→ Critical section

# Condition Variables

- Some problems will require “waiting for another thread” – **Synchronization**
  - ◆ Ex. Threads for I/O processing
  - ◆ Other synchronization problems: Bounded-Buffer, Producer-Consumer, Readers-Writers, Dining Philosophers Problem