

# CMSC 125: Operating Systems

- ❑ Instructor: **Joseph Anthony C. Hermocilla**
- ❑ Email: [jchermocilla@up.edu.ph](mailto:jchermocilla@up.edu.ph)
- ❑ Web: <https://jachermocilla.org>



# Resources

Book: <https://pages.cs.wisc.edu/~remzi/OSTEP/>

Slides Template:

<https://pages.cs.wisc.edu/~remzi/OSTEP/Educators-Slides/Youjip/>



# Acknowledgement

- ▣ This lecture slide set was initially developed for Operating System course in Computer Science Dept. at Hanyang University. This lecture slide set is for OSTEP book written by Remzi and Andrea at University of Wisconsin.

# **30. Condition Variables**

**Operating System: Three Easy Pieces**

# Condition Variables

- There are many cases where a thread wishes to check whether a **condition** is true before continuing its execution
- Example:
  - ◆ A parent thread might wish to check whether a child thread has *completed*
  - ◆ This is often called a `join()`

# Condition Variables (Cont.)

## A Parent Waiting For Its Child

```
1      void *child(void *arg) {
2          printf("child\n");
3          // XXX how to indicate we are done?
4          return NULL;
5      }
6
7      int main(int argc, char *argv[]) {
8          printf("parent: begin\n");
9          pthread_t c;
10         Pthread_create(&c, NULL, child, NULL); // create child
11         // XXX how to wait for child?
12         printf("parent: end\n");
13         return 0;
14     }
```

What we would like to see here is:

```
parent: begin
child
parent: end
```

# Parent waiting fore child: Spin-based Approach

```
1      volatile int done = 0;
2
3      void *child(void *arg) {
4          printf("child\n");
5          done = 1;
6          return NULL;
7      }
8
9      int main(int argc, char *argv[]) {
10         printf("parent: begin\n");
11         pthread_t c;
12         Pthread_create(&c, NULL, child, NULL); // create child
13         while (done == 0)
14             ; // spin
15         printf("parent: end\n");
16         return 0;
17     }
```

- ◆ This is hugely inefficient as the parent spins and **wastes CPU time**.

# How to wait for a condition

## □ What is a Condition Variable?

- ◆ An **explicit queue** that threads can put themselves on when some state of execution( the **condition**) is not as desired (by waiting on the condition)
  - **Waiting** on the condition
    - A threads put itself in the queue until another thread signals, ensures that thread does not spin
  - **Signaling** on the condition
    - Some other thread, *when it changes said state(condition)*, can wake one of those waiting threads and allow them to continue



# Definition and Routines

## ❑ Declare condition variable

```
pthread_cond_t c;
```

- ◆ Proper initialization is required

## ❑ Operation (the POSIX calls)

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);    // wait()  
pthread_cond_signal(pthread_cond_t *c);                      // signal()
```

- ◆ The wait() call takes a mutex as a parameter (assumed to be locked) because
  - The wait() call releases the lock and puts the calling thread to sleep
  - When the thread wakes up, it must re-acquire the lock

# Parent waiting for Child: Use a condition variable

```
1      int done = 0;
2      pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3      pthread_cond_t c = PTHREAD_COND_INITIALIZER;
4
5      void thr_exit() {
6          Pthread_mutex_lock(&m);
7          done = 1;
8          Pthread_cond_signal(&c);
9          Pthread_mutex_unlock(&m);
10     }
11
12     void *child(void *arg) {
13         printf("child\n");
14         thr_exit();
15         return NULL;
16     }
17
18     void thr_join() {
19         Pthread_mutex_lock(&m);
20         while (done == 0)
21             Pthread_cond_wait(&c, &m);
22         Pthread_mutex_unlock(&m);
23     }
24
```

## Parent waiting for Child: Use a condition variable

(cont.)

```
25     int main(int argc, char *argv[]) {
26         printf("parent: begin\n");
27         pthread_t p;
28         Pthread_create(&p, NULL, child, NULL);
29         thr_join();
30         printf("parent: end\n");
31         return 0;
32     }
```

# Parent waiting for Child: Use a condition variable

## □ Parent:

- ◆ Create the child thread and continues running itself
- ◆ Call into `thr_join()` to wait for the child thread to complete
  - Acquire the lock
  - Check if the child is done
  - Put itself to sleep by calling `wait()`
  - Release the lock

## □ Child:

- ◆ Print the message “child”
- ◆ Call `thr_exit()` to wake the parent thread
  - Grab the lock
  - Set the state variable `done`
  - Signal the parent thus waking it

# The importance of the state variable `done`

```
1      void thr_exit() {  
2          pthread_mutex_lock(&m);  
3          pthread_cond_signal(&c);  
4          pthread_mutex_unlock(&m);  
5      }  
6  
7      void thr_join() {  
8          pthread_mutex_lock(&m);  
9          pthread_cond_wait(&c, &m);  
10         pthread_mutex_unlock(&m);  
11     }
```

`thr_exit()` and `thr_join()` without variable `done`

- ◆ Imagine the case where the *child runs immediately*
  - The child will signal, but there is no thread asleep on the condition
  - When the parent runs, it will call wait and be **stuck**
  - No thread will ever wake it

# Another poor implementation

```
1      void thr_exit() {
2          done = 1;
3          Pthread_cond_signal(&c);
4      }
5
6      void thr_join() {
7          if (done == 0)
8              Pthread_cond_wait(&c);
9      }
```

- ◆ The issue here is a subtle **race condition**
  - The parent calls `thr_join()`
    - The parent checks the value of `done`
    - It will see that it is 0 and try to go to sleep
    - *Just before* it calls wait to go to sleep, the parent is interrupted and the child runs
  - The child changes the state variable `done` to 1 and signals
    - But no thread is waiting and thus no thread is woken
    - When the parent runs again, it sleeps forever

# The Producer / Consumer (Bounded Buffer) Problem

## □ **Producer**

- ◆ **Produce** data items
- ◆ Wish to place data items in a buffer

## □ **Consumer**

- ◆ Grab data items out of the buffer **consume** them in some way

## □ **Example: Multi-threaded web server**

- ◆ *A producer* puts HTTP requests into a work queue
- ◆ *Consumer threads* take requests out of this queue and process them

# Bounded Buffer

- A bounded buffer is used when you pipe the output of one program into another
  - ◆ **Example:** `grep foo file.txt | wc -l`
    - The `grep` process is the producer
    - The `wc` process is the consumer
    - Between them is an in-kernel bounded buffer (a pipe)
  - ◆ Bounded buffer is shared resource → **Synchronized access** is required



# The Put and Get Routines (Version 1)

```
1      int buffer;
2      int count = 0;    // initially, empty
3
4      void put(int value) {
5          assert(count == 0);
6          count = 1;
7          buffer = value;
8      }
9
10     int get() {
11         assert(count == 1);
12         count = 0;
13         return buffer;
14     }
```

- ◆ Only put data into the buffer when `count` is zero
  - i.e., when the buffer is *empty*
- ◆ Only get data from the buffer when `count` is one
  - i.e., when the buffer is *full*

# Producer/Consumer Threads (Version 1)

```
1      void *producer(void *arg) {
2          int i;
3          int loops = (int) arg;
4          for (i = 0; i < loops; i++) {
5              put(i);
6          }
7      }
8
9      void *consumer(void *arg) {
10         int i;
11         while (1) {
12             int tmp = get();
13             printf("%d\n", tmp);
14         }
15     }
```

- ◆ **Producer** puts an integer into the shared buffer `loops` number of times
- ◆ **Consumer** gets the data out of that shared buffer

# Producer/Consumer: Single CV and If Statement

- A single condition variable `cond` and associated lock `mutex`

```
1      int loops; cond_t cond;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              if (count == 1)                       // p2
9                  Pthread_cond_wait(&cond, &mutex); // p3
10             put(i);                                // p4
11             Pthread_cond_signal(&cond);            // p5
12             Pthread_mutex_unlock(&mutex);          // p6
13         }
14     }
15
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);           // c1
```

# Producer/Consumer: Single CV and If Statement

```
20         if (count == 0)                                // c2
21             Pthread_cond_wait(&cond, &mutex);           // c3
22         int tmp = get();                                  // c4
23         Pthread_cond_signal(&cond);                      // c5
24         Pthread_mutex_unlock(&mutex);                    // c6
25         printf("%d\n", tmp);
26     }
27 }
```


- ♦ p1-p3: A producer waits for the buffer to be empty
- ♦ c1-c3: A consumer waits for the buffer to be full
- ♦ With just *a single producer* and *a single consumer*, the code works

If we have **more than** one of producer and consumer?

# Thread Trace: Broken Solution (Version 1)

$T_{c1}$	State	$T_{c2}$	State	$T_p$	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep		Ready	p1	Running	0	
	Sleep		Ready	p2	Running	0	
	Sleep		Ready	p4	Running	1	Buffer now full
	Ready		Ready	p5	Running	1	$T_{c1}$ awoken
	Ready		Ready	p6	Running	1	
	Ready		Ready	p1	Running	1	
	Ready		Ready	p2	Running	1	
	Ready		Ready	p3	Sleep	1	Buffer full; sleep
	Ready	c1	Running		Sleep	1	$T_{c2}$ sneaks in ...
	Ready	c2	Running		Sleep	1	
	Ready	c4	Running		Sleep	0	... and grabs data
	Ready	c5	Running		Ready	0	$T_p$ awoken
	Ready	c6	Running		Ready	0	
c4	Running		Ready		Ready	0	Oh oh! No data

c3 is skipped since  
count is 1



# Thread Trace: Broken Solution (Version 1)

- The problem arises for a simple reason:
  - ◆ After the producer woke  $T_{c1}$ , but before  $T_{c1}$  ever ran, the state of the bounded buffer *changed by*  $T_{c2}$
  - ◆ There is no guarantee that when the woken thread runs, the state will still be as desired → Mesa semantics
    - Virtually every system ever built employs *Mesa semantics*
  - ◆ Hoare semantics provides a stronger guarantee that the woken thread will run immediately upon being woken

# Producer/Consumer: Single CV and While

- Consumer  $T_{c1}$  wakes up and **re-checks** the state of the shared variable
  - ◆ If the buffer is empty, the consumer simply goes back to sleep

```
1      cond_t cond;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              while (count == 1)                    // p2
9                  Pthread_cond_wait(&cond, &mutex); // p3
10             put(i);                                // p4
11             Pthread_cond_signal(&cond);            // p5
12             Pthread_mutex_unlock(&mutex);          // p6
13         }
14     }
15
```

# Producer/Consumer: Single CV and While

```
(Cont.)
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);           // c1
20             while (count == 0)                    // c2
21                 Pthread_cond_wait(&cond, &mutex); // c3
22             int tmp = get();                       // c4
23             Pthread_cond_signal(&cond);           // c5
24             Pthread_mutex_unlock(&mutex);         // c6
25             printf("%d\n", tmp);
26         }
27     }
```

- ◆ A simple rule to remember with condition variables is to **always use while loops**
- ◆ However, this code still has a bug (*next page*)



# Thread Trace: Broken Solution (Version 2)

$T_{c1}$	State	$T_{c2}$	State	$T_p$	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep	c1	Running		Ready	0	
	Sleep	c2	Running		Ready	0	
	Sleep	c3	Sleep		Ready	0	Nothing to get
	Sleep		Sleep	p1	Running	0	
	Sleep		Sleep	p2	Running	0	
	Sleep		Sleep	p4	Running	1	Buffer now full
	Ready		Sleep	p5	Running	1	$T_{c1}$ awoken
	Ready		Sleep	p6	Running	1	
	Ready		Sleep	p1	Running	1	
	Ready		Sleep	p2	Running	1	
	Ready		Sleep	p3	Sleep	1	Must sleep (full)
c2	Running		Sleep		Sleep	1	Recheck condition
c4	Running		Sleep		Sleep	0	$T_{c1}$ grabs data
c5	Running		Ready		Sleep	0	<b>Oops! Woke <math>T_{c2}</math></b>

## Thread Trace: Broken Solution (Version 2) (Cont.)

$T_{c1}$	State	$T_{c2}$	State	$T_p$	State	Count	Comment
...	...	...	...	...	...	...	(cont.)
c6	Running		Ready		Sleep	0	
c1	Running		Ready		Sleep	0	
c2	Running		Ready		Sleep	0	
c3	Sleep		Ready		Sleep	0	Nothing to get
	Sleep	c2	Running		Sleep	0	
	Sleep	c3	Sleep		Sleep	0	<b>Everyone asleep ...</b>

- ◆ A consumer should not wake other consumers, only producers, and vice-versa

# The single Buffer Producer/Consumer Solution

- Use **two** condition variables and while
  - ◆ **Producer** threads wait on the condition `empty`, and signals `fill`
  - ◆ **Consumer** threads wait on `fill` and signal `empty`

```
1      cond_t empty, fill;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);
8              while (count == 1)
9                  Pthread_cond_wait(&empty, &mutex);
10             put(i);
11             Pthread_cond_signal(&fill);
12             Pthread_mutex_unlock(&mutex);
13         }
14     }
15
```

# The single Buffer Producer/Consumer Solution

```
(Cont.)
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);
20             while (count == 0)
21                 Pthread_cond_wait(&fill, &mutex);
22             int tmp = get();
23             Pthread_cond_signal(&empty);
24             Pthread_mutex_unlock(&mutex);
25             printf("%d\n", tmp);
26         }
27     }
```

# The Final Producer/Consumer Solution

- More **concurrency** and **efficiency** → Add more buffer slots
  - ◆ Allow concurrent production or consuming to take place
  - ◆ Reduce context switches

```
1      int buffer[MAX];
2      int fill = 0;
3      int use = 0;
4      int count = 0;
5
6      void put(int value) {
7          buffer[fill] = value;
8          fill = (fill + 1) % MAX;
9          count++;
10     }
11
12     int get() {
13         int tmp = buffer[use];
14         use = (use + 1) % MAX;
15         count--;
16         return tmp;
17     }
```

The Final Put and Get Routines

# The Final Producer/Consumer Solution (Cont.)

```
1      cond_t empty, fill;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              while (count == MAX)                 // p2
9                  Pthread_cond_wait(&empty, &mutex); // p3
10             put(i);                               // p4
11             Pthread_cond_signal(&fill);           // p5
12             Pthread_mutex_unlock(&mutex);         // p6
13         }
14     }
15
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);           // c1
20             while (count == 0)                     // c2
21                 Pthread_cond_wait(&fill, &mutex); // c3
22             int tmp = get();                       // c4
```

# The Final Producer/Consumer Solution (Cont.)

```
(Cont.)  
23         Pthread_cond_signal(&empty);           // c5  
24         Pthread_mutex_unlock(&mutex);          // c6  
25         printf("%d\n", tmp);  
26     }  
27 }
```

## The Final Working Solution (Cont.)

- ♦ **p2: A producer** only sleeps if all buffers are currently filled
- ♦ **c2: A consumer** only sleeps if all buffers are currently empty

# Covering Conditions

- Assume there are zero bytes free
  - ◆ Thread  $T_a$  calls `allocate(100)`.
  - ◆ Thread  $T_b$  calls `allocate(10)`.
  - ◆ Both  $T_a$  and  $T_b$  wait on the condition and go to sleep.
  - ◆ Thread  $T_c$  calls `free(50)`.

Which waiting thread should be woken up?



## Covering Conditions (Cont.)

```
1      // how many bytes of the heap are free?
2      int bytesLeft = MAX_HEAP_SIZE;
3
4      // need lock and condition too
5      cond_t c;
6      mutex_t m;
7
8      void *
9      allocate(int size) {
10         Pthread_mutex_lock(&m);
11         while (bytesLeft < size)
12             Pthread_cond_wait(&c, &m);
13         void *ptr = ...;           // get mem from heap
14         bytesLeft -= size;
15         Pthread_mutex_unlock(&m);
16         return ptr;
17     }
18
19     void free(void *ptr, int size) {
20         Pthread_mutex_lock(&m);
21         bytesLeft += size;
22         Pthread_cond_signal(&c);   // whom to signal??
23         Pthread_mutex_unlock(&m);
24     }
```

## Covering Conditions (Cont.)

### ▣ Solution (Suggested by Lampson and Redell)

- ◆ **Replace** `pthread_cond_signal()` **with** `pthread_cond_broadcast()`
- ◆ `pthread_cond_broadcast()`
  - Wake up **all waiting threads**
  - Cost: too many threads might be woken
  - Threads that shouldn't be awake will simply wake up, re-check the condition, and then go back to sleep